

Q0 Before you start; before you submit

I did not collaborate with anyone on this homework.

Q1 Self assessment

The part of HW3 with the most room for improvement is Question 2(c), the lower bound on the runtime of the recursive COUNT algorithm. Here is what I submitted (after the pseudocode):

Let $T(n)$ be its running time. For $n \geq 2$, $T(n) \geq T(n-1) + T(n-2)$. This recurrence is bounded below by a constant multiple of the Fibonacci numbers, so $T(n) = \Omega(\varphi^n)$, where $\varphi = (1 + \sqrt{5})/2 > 1$. The runtime is therefore exponential in n .

What was deficient: the question asked me to *show* a lower bound, but I only stated one. The sentence “this recurrence is bounded below by a constant multiple of the Fibonacci numbers” is the whole argument, and I never proved it. Going from $T(n) \geq T(n-1) + T(n-2)$ to $T(n) = \Omega(\varphi^n)$ needs an induction (or an unrolling of the recurrence), and I skipped that step. So the conclusion is true, but the write-up does not justify it.

How to fix it: the posted solution does something simpler and actually shows the work. Since T is nondecreasing, $T(n-1) \geq T(n-2)$, so $T(n) \geq 2T(n-2)$. Unrolling this k times gives $T(n) \geq 2^k T(n-2k)$, and taking $k = n/2$ gives $T(n) \geq 2^{n/2} T(0) = \Omega(2^{n/2})$. That bound is weaker than φ^n , but it is exponential in n , which is all the question asked for, and every step is written down. If I wanted to keep the φ^n bound, I would need to prove $T(n) \geq c\varphi^n$ by induction using $\varphi^2 = \varphi + 1$.

Q2.1 Fixed-color-sequence paths (a)

For $0 \leq i \leq k$ and a vertex $v \in V$, define

$$R(i, v) = \begin{cases} \text{true} & \text{if there is a path of length } i \text{ from } s \text{ to } v \text{ with color sequence } (c_1, \dots, c_i), \\ \text{false} & \text{otherwise.} \end{cases}$$

(A path of length 0 is just a single vertex with no edges, and its color sequence is empty.) The answer to the whole problem is $R(k, t)$.

Base case ($i = 0$). The only path of length 0 starting at s is the path that stays at s . So

$$R(0, v) = \begin{cases} \text{true} & \text{if } v = s, \\ \text{false} & \text{if } v \neq s. \end{cases}$$

Recursive case ($i \geq 1$).

$$R(i, v) = \bigvee_{\substack{u: (u, v) \in E \\ \text{color}(u, v) = c_i}} R(i-1, u).$$

In words: $R(i, v)$ is true exactly when some neighbor u of v , joined to v by an edge of color c_i , has $R(i-1, u)$ true. (An OR over zero terms is false, so if v has no incident edge of color c_i , then $R(i, v)$ is false.)

Justification. I show both directions of the “exactly when”.

Suppose $R(i, v)$ is true, so there is a path $P = (e_1, \dots, e_i)$ from s to v with colors $(c_1, \dots, c_i)$. Its last edge $e_i = (u, v)$ has color c_i for some neighbor u of v . Removing e_i leaves the path $(e_1, \dots, e_{i-1})$, which goes from s to u , has length $i-1$, and has color sequence $(c_1, \dots, c_{i-1})$. So $R(i-1, u)$ is true, and u is one of the terms in the OR. Hence the right-hand side is true.

Conversely, suppose the right-hand side is true. Then there is a neighbor u of v with $\text{color}(u, v) = c_i$ and $R(i-1, u)$ true. Take a path of length $i-1$ from s to u with colors $(c_1, \dots, c_{i-1})$ and append the edge (u, v) . This gives a path of length i from s to v with colors $(c_1, \dots, c_i)$. Paths do not need to be simple, so repeating a vertex or an edge is fine. Hence $R(i, v)$ is true.

Q2.2 Fixed-color-sequence paths (b)

I assume each adjacency-list entry for a neighbor u of v also stores the color of the edge (u, v) , so looking up a color takes $O(1)$ time. (If colors are stored separately, e.g. in a hash table keyed by edge, the lookup is still $O(1)$.)

Input: Adjacency lists of $G = (V, E)$ with edge colors; vertices s, t ; colors $c_1, \dots, c_k$

Output: true if there is a path of length k from s to t with color sequence $(c_1, \dots, c_k)$, else false

```

1: function COLORPATH( $G, s, t, (c_1, \dots, c_k)$ )
2:   initialize a table  $R[0 \dots k][V]$  of Booleans
3:   for each  $v \in V$  do                                     ▷ base case
4:      $R[0][v] \leftarrow \text{false}$ 
5:   end for
6:    $R[0][s] \leftarrow \text{true}$ 
7:   for  $i = 1, \dots, k$  do                                   ▷ fill layer  $i$  from layer  $i - 1$ 
8:     for each  $v \in V$  do
9:        $R[i][v] \leftarrow \text{false}$ 
10:      for each neighbor  $u$  in the adjacency list of  $v$  do
11:        if  $\text{color}(u, v) = c_i$  and  $R[i - 1][u] = \text{true}$  then
12:           $R[i][v] \leftarrow \text{true}$ 
13:        end if
14:      end for
15:    end for
16:  end for
17:  return  $R[k][t]$ 
18: end function

```

Order of evaluation. Layer i of the table only uses entries from layer $i - 1$, and the loop fills the layers in the order $0, 1, \dots, k$. So every value is computed before it is needed, and the table entry $R[i][v]$ holds $R(i, v)$ from part (a). The answer is in $R[k][t]$.

Running time. Filling layer 0 takes $O(n)$ time. For one value of i , the loop over v walks the entire adjacency list of every vertex. Each edge (u, v) appears twice over all the lists (once for u , once for v), so one layer does $O(n + 2m) = O(n + m)$ work, with $O(1)$ work per list entry. There are k layers, so the total is $O(n + k(n + m))$. Since $m \geq n$, this is $O(km)$. (The table uses $O(kn)$ space; keeping only layers $i - 1$ and i would reduce this to $O(n)$, since the answer only needs the last layer.)

Q3.1 Back to college (a)

Idea. Starting from the current town with a full pack, drive as far as the pack allows, i.e. keep going while the next distance still fits in the remaining fuel. If the car gets to town n , we are done. Otherwise, refuel at the farthest town that was reached, and repeat from there. So the greedy rule is: “always refuel as late as possible.”

For the pseudocode I use a helper notation: for $0 \leq a < b \leq n$, let

$$\text{dist}(a, b) = d_{a+1} + d_{a+2} + \cdots + d_b$$

be the distance from town a to town b . A set S of refueling towns is *valid* if every stretch between consecutive elements of $\{0\} \cup S \cup \{n\}$ has distance at most R .

Input: $R > 0$ and distances $D[1 \dots n]$ with each $D[i] \leq R$

Output: A smallest valid set S of refueling towns

```

1: function GREEDYREFUEL( $R, D[1 \dots n]$ )
2:    $S \leftarrow \emptyset$ 
3:    $i \leftarrow 0$                                 ▷ current town; the car has a full pack here
4:   while  $i < n$  do
5:      $j \leftarrow i, \text{fuel} \leftarrow R$ 
6:     while  $j < n$  and  $\text{fuel} \geq D[j + 1]$  do      ▷ drive as far as this pack allows
7:        $\text{fuel} \leftarrow \text{fuel} - D[j + 1]$ 
8:        $j \leftarrow j + 1$ 
9:     end while
10:    if  $j = n$  then
11:      return  $S$                                 ▷ reached the last town, no more refueling needed
12:    end if
13:    add  $j$  to  $S$                                 ▷ refuel at the farthest town reached
14:     $i \leftarrow j$ 
15:  end while
16:  return  $S$                                     ▷ only reached when  $n = 0$ 
17: end function

```

Why it always makes progress and outputs a valid set. Each time the inner loop starts, $\text{fuel} = R \geq D[i + 1]$ because every $d_i \leq R$. So the inner loop runs at least once and j ends strictly larger than i . That means i strictly increases in every iteration of the outer loop, so the algorithm terminates. Also, by the inner loop condition, the distance from i to j is at most R , so each stretch between consecutive stops (and from the last stop to town n) fits in one pack. Hence the output is a valid set.

Running time. The index j starts at the old value of i and is only ever incremented, and the next outer iteration starts j from where the last one ended. So over the whole run, the inner loop body executes at most n times in total. Every other operation is $O(1)$ per iteration. The total running time is $O(n)$. The extra space is $O(1)$ besides the output set.

Q3.2 Back to college (b)

Throughout, I use $\text{dist}(a, b) = d_{a+1} + \dots + d_b$ from part (a), and the following two facts.

- **Fact 1 (sub-stretches).** If $a \leq a' < b' \leq b$ then $\text{dist}(a', b') \leq \text{dist}(a, b)$, because all d_i are nonnegative. So if a stretch fits in one pack, so does any stretch inside it.
- **Fact 2 (what the greedy choice means).** Let $a = i_{\text{diff}-1}$, or $a = 0$ if $i_{\text{diff}} = i_1$. The algorithm chose i_{diff} by starting at a with a full pack and driving as far as possible. So $\text{dist}(a, i_{\text{diff}}) \leq R$, and since the algorithm did not stop at town n , we have $i_{\text{diff}} < n$ and $\text{dist}(a, i_{\text{diff}} + 1) > R$.

Which town j to exchange. Let j be the smallest town in OPT that is larger than a . I claim j exists and $a < j < i_{\text{diff}}$.

Existence. a is a refueling stop of OPT (or $a = 0$, the start), so in OPT the car leaves a with a full pack. If OPT had no stop in $\{a + 1, \dots, i_{\text{diff}}\}$, then the car would have to drive from a to some town $\geq i_{\text{diff}} + 1$ (either the next stop or town n) on one pack. By Fact 1 that stretch has distance at least $\text{dist}(a, i_{\text{diff}} + 1) > R$, which contradicts the validity of OPT. So OPT has a stop in $\{a + 1, \dots, i_{\text{diff}}\}$, and hence j exists with $a < j \leq i_{\text{diff}}$.

Strictness. $i_{\text{diff}} \notin \text{OPT}$ but $j \in \text{OPT}$, so $j \neq i_{\text{diff}}$, hence $j < i_{\text{diff}}$.

Now define

$$\text{OPT}' = (\text{OPT} \setminus \{j\}) \cup \{i_{\text{diff}}\}.$$

OPT' contains $i_1, \dots, i_{\text{diff}}$. The towns $i_1, \dots, i_{\text{diff}-1}$ are in OPT and are all $\leq a < j$, so none of them is the removed town j . They stay in OPT', and i_{diff} was added.

OPT' is valid. Let $p < q$ be two consecutive elements of $\{0\} \cup \text{OPT}' \cup \{n\}$; I show $\text{dist}(p, q) \leq R$. Note that $i_{\text{diff}} \in \text{OPT}'$, so it is impossible that $p < i_{\text{diff}} < q$. That leaves two cases.

Case $q \leq i_{\text{diff}}$. The town a is in $\{0\} \cup \text{OPT}'$ (it was in OPT and $a \neq j$), and $a < j < i_{\text{diff}}$, so $a < q$. Since p is the element right before q , we get $p \geq a$. So $a \leq p < q \leq i_{\text{diff}}$, and by Facts 1 and 2, $\text{dist}(p, q) \leq \text{dist}(a, i_{\text{diff}}) \leq R$.

Case $p \geq i_{\text{diff}}$. Above i_{diff} , the sets OPT and OPT' agree, because the only changes were at j and i_{diff} , both $\leq i_{\text{diff}}$.

- If $p > i_{\text{diff}}$, then p and q are also consecutive in $\{0\} \cup \text{OPT} \cup \{n\}$, so $\text{dist}(p, q) \leq R$ by the validity of OPT.
- If $p = i_{\text{diff}}$, then q is the smallest element of $\text{OPT}' \cup \{n\}$ above i_{diff} , and since OPT and OPT' agree there, q is also the smallest element of $\text{OPT} \cup \{n\}$ above i_{diff} . Let p' be the element of $\{0\} \cup \text{OPT}$ right before q . Then $p' \neq i_{\text{diff}}$ (as $i_{\text{diff}} \notin \text{OPT}$) and there is no element of OPT strictly between i_{diff} and q , so $p' < i_{\text{diff}}$. OPT is valid, so $\text{dist}(p', q) \leq R$. Since $p' < i_{\text{diff}} < q$, Fact 1 gives $\text{dist}(i_{\text{diff}}, q) \leq \text{dist}(p', q) \leq R$.

In all cases every stretch of OPT' has distance at most R , so OPT' is a valid solution.

OPT' is optimal. We removed one town ($j \in \text{OPT}$) and added one town ($i_{\text{diff}} \notin \text{OPT}$), so $|\text{OPT}'| = |\text{OPT}|$. OPT' is valid and has the same size as an optimal solution, so OPT' is optimal.

Q3.3 Back to college (c)

Part (b) says: if an optimal solution agrees with the greedy choices $i_1, \dots, i_{\text{diff}-1}$ but not with i_{diff} , we can fix it into another optimal solution that also agrees with i_{diff} . The proof only used that OPT is optimal and contains $i_1, \dots, i_{\text{diff}-1}$, so it works for any such OPT and any position diff .

So we can repeat it. Start with any optimal solution OPT . If it does not contain all of $i_1, \dots, i_k$, apply part (b) at the first disagreement to get an optimal solution that agrees with one more greedy town. Each application strictly increases the number of leading greedy towns that are contained, and there are only k of them, so after at most k rounds we get an optimal solution OPT^* with $\{i_1, \dots, i_k\} \subseteq \text{OPT}^*$. (Formally this is an induction on t : for every $t \leq k$ there is an optimal solution containing $i_1, \dots, i_t$.)

Now compare sizes. ALG is a valid solution (shown in part (a)), so $|\text{ALG}| \geq |\text{OPT}^*|$ because OPT^* is a smallest valid solution. On the other hand $\text{ALG} \subseteq \text{OPT}^*$, so $|\text{ALG}| \leq |\text{OPT}^*|$. Hence $|\text{ALG}| = |\text{OPT}^*|$, so ALG is optimal, i.e., the greedy algorithm always outputs a smallest valid set of refueling towns.

Q3.4 Back to college (d)

The claim is false.

Counterexample. Take $R = 2$, $n = 3$, distances $D = [1, 1, 1]$, and costs $C = [1, 10, 1]$.

Greedy output. From town 0 with a full pack the car can drive $1 + 1 = 2 \leq R$ to town 2, but not $1 + 1 + 1 = 3 > R$ to town 3. So the greedy algorithm refuels at town 2, then drives $1 \leq R$ to town 3. Its output is $S = \{2\}$ with total cost $c_2 = 10$.

Better solution. $S' = \{1\}$ is valid: the stretch from 0 to 1 has distance $1 \leq 2$ and the stretch from 1 to 3 has distance $1 + 1 = 2 \leq 2$. Its cost is $c_1 = 1 < 10$. So the greedy output is not optimal.

Where the proof fails. In part (b), OPT' was obtained from OPT by swapping j for i_{diff} , and optimality followed from $|\text{OPT}'| = |\text{OPT}|$. With costs, the swap changes the objective by $c_{i_{\text{diff}}} - c_j$, which can be positive (here $c_2 - c_1 = 9$). So OPT' is still valid, but it can be strictly worse than OPT , and the exchange step no longer produces an optimal solution.

Q4.1 Dynamic shortest paths (a)

Conventions: $D(y, z) = \infty$ if there is no path from y to z , $D(y, y) = 0$, and the minimum over an empty set is ∞ . I also use the standard fact that in a graph with no negative-weight cycle, for every u, v with a path from u to v there is a shortest path that is *simple*: given any walk from u to v , repeatedly cut out a cycle; this never increases the length because every cycle has weight ≥ 0 .

Expressions. For every $u \in V$,

$$D_{new}(v_{new}, u) = \min_{x: (v_{new}, x) \in E_{new}} (w(v_{new}, x) + D(x, u)),$$

$$D_{new}(u, v_{new}) = \min_{y: (y, v_{new}) \in E_{new}} (D(u, y) + w(y, v_{new})).$$

(Also $D_{new}(v_{new}, v_{new}) = 0$.)

Correctness of the first expression. Fix $u \in V$ and let M be the right-hand side.

$D_{new}(v_{new}, u) \leq M$. If $M = \infty$ there is nothing to show. Otherwise take an out-edge (v_{new}, x) attaining the minimum, so $D(x, u) < \infty$. Take a shortest path from x to u in G (length $D(x, u)$) and put the edge (v_{new}, x) in front of it. This is a walk from v_{new} to u in G_{new} of length $w(v_{new}, x) + D(x, u) = M$. By the fact above, there is a path of length at most M , so $D_{new}(v_{new}, u) \leq M$.

$D_{new}(v_{new}, u) \geq M$. If $D_{new}(v_{new}, u) = \infty$ there is nothing to show. Otherwise take a simple shortest path P from v_{new} to u in G_{new} . Since $u \neq v_{new}$, P has at least one edge, and its first edge is some $(v_{new}, x) \in E_{new}$. Since P is simple, v_{new} does not appear again, so the rest of P is a path from x to u that uses only vertices of V and edges of E . Its length is at least $D(x, u)$. Hence

$$D_{new}(v_{new}, u) = \text{len}(P) \geq w(v_{new}, x) + D(x, u) \geq M.$$

This is the same idea as the Bellman-Ford relaxation step: a shortest path to u is an edge out of the source followed by a shortest path from the other endpoint.

Correctness of the second expression. This is symmetric. A simple shortest path from u to v_{new} in G_{new} has v_{new} only as its last vertex, so its last edge is some $(y, v_{new}) \in E_{new}$ and the part before it is a path from u to y inside G , of length at least $D(u, y)$. That gives $D_{new}(u, v_{new}) \geq \min_y (D(u, y) + w(y, v_{new}))$. Conversely, a shortest u -to- y path in G followed by the edge (y, v_{new}) is a walk of length $D(u, y) + w(y, v_{new})$, and it can be shortened to a path, so $D_{new}(u, v_{new}) \leq \min_y (D(u, y) + w(y, v_{new}))$.

Evaluation time. v_{new} has at most n out-edges and at most n in-edges, so each minimum ranges over at most n terms. Each term is a table lookup plus an addition, which is $O(1)$. So each of the two values takes $O(n)$ time for a fixed u , and computing them for all n vertices u takes $O(n^2)$ time in total.

Q4.2 Dynamic shortest paths (b)

Expression. For all $u, v \in V$,

$$D_{new}(u, v) = \min \{ D(u, v), D_{new}(u, v_{new}) + D_{new}(v_{new}, v) \},$$

where the two D_{new} terms on the right are the values computed in part (a).

This is exactly one round of Floyd-Warshall with v_{new} as the last “intermediate” vertex: a shortest path either does not go through v_{new} , or it does, and then it splits into a shortest path to v_{new} and a shortest path from v_{new} to v .

Correctness. Fix $u, v \in V$ and let M be the right-hand side.

$D_{new}(u, v) \leq M$. First, every path in G is also a path in G_{new} , so $D_{new}(u, v) \leq D(u, v)$. Second, if both $D_{new}(u, v_{new})$ and $D_{new}(v_{new}, v)$ are finite, then a shortest path from u to v_{new} followed by a shortest path from v_{new} to v is a walk from u to v in G_{new} of length $D_{new}(u, v_{new}) + D_{new}(v_{new}, v)$. Since G_{new} has no negative cycle, this walk can be shortened to a path with no larger length, so $D_{new}(u, v) \leq D_{new}(u, v_{new}) + D_{new}(v_{new}, v)$. Together, $D_{new}(u, v) \leq M$.

$D_{new}(u, v) \geq M$. If $D_{new}(u, v) = \infty$ we are done. Otherwise take a simple shortest path P from u to v in G_{new} .

- If P does not visit v_{new} , then P only uses vertices of V and edges of E , so it is a path in G and $\text{len}(P) \geq D(u, v) \geq M$.
- If P visits v_{new} , it does so exactly once because P is simple. Split P at v_{new} into P_1 (from u to v_{new}) and P_2 (from v_{new} to v). Both are paths in G_{new} , so $\text{len}(P_1) \geq D_{new}(u, v_{new})$ and $\text{len}(P_2) \geq D_{new}(v_{new}, v)$. Hence $\text{len}(P) = \text{len}(P_1) + \text{len}(P_2) \geq D_{new}(u, v_{new}) + D_{new}(v_{new}, v) \geq M$.

In both cases $D_{new}(u, v) = \text{len}(P) \geq M$.

Evaluation time. After part (a), the values $D_{new}(u, v_{new})$ and $D_{new}(v_{new}, v)$ are stored for every u and v . So for each pair (u, v) the expression is three table lookups, one addition and one minimum, i.e. $O(1)$ time. There are n^2 pairs, so the total is $O(n^2)$. Together with the $O(n^2)$ time of part (a), the whole table D_{new} (including the row and column of v_{new}) is computed in $O(n^2)$ time.

Q5.1 MST mashup (a)

The first two sentences are fine: with distinct weights, Kruskal’s algorithm is deterministic and has a single possible output.

The third sentence does not follow from the second. “Kruskal’s algorithm is correct” means that its output is *some* minimum spanning tree. It does not say that *every* minimum spanning tree is an output of Kruskal’s algorithm. So “the algorithm has one output” plus “that output is an MST” does not rule out other MSTs that the algorithm simply never produces.

Here is why this kind of reasoning is invalid in general. Take a triangle whose three edges all have weight 1, and run Kruskal’s algorithm with a fixed tie-breaking rule (say, by edge index). Then the algorithm is deterministic and has exactly one output, and it is correct, so its output is an MST. But the triangle has three different MSTs (drop any one edge). So “deterministic and correct” by itself never implies uniqueness of the optimum. The claim happens to be true for distinct weights, but it needs a real argument about the weights, which is part (b).

Q5.2 MST mashup (b)

Let G be a connected graph with distinct edge weights, and suppose for contradiction that $T_1 \neq T_2$ are both minimum spanning trees of G .

Choosing the edge to exchange. Since $T_1 \neq T_2$ and both are sets of edges, the symmetric difference $T_1 \Delta T_2 = (T_1 \setminus T_2) \cup (T_2 \setminus T_1)$ is nonempty. Let e be the edge of minimum weight in $T_1 \Delta T_2$; it is unique since weights are distinct. By symmetry (swap the names of T_1 and T_2 if needed), assume $e \in T_1 \setminus T_2$.

Finding a heavier edge f in T_2 . T_2 is a spanning tree, so its two endpoints of e are connected by a path in T_2 . Adding e to T_2 therefore creates a cycle C in $T_2 \cup \{e\}$, and $e \in C$. Not every edge of C can be in T_1 : otherwise C would be a cycle inside the tree T_1 , which is impossible. So C has some edge $f \notin T_1$. Since $f \neq e$ (as $e \in T_1$) and every edge of C other than e lies in T_2 , we get $f \in T_2 \setminus T_1 \subseteq T_1 \Delta T_2$. By the choice of e as the unique lightest edge of $T_1 \Delta T_2$, and since $f \neq e$,

$$w(f) > w(e).$$

The exchange. Let $T'_2 = (T_2 \setminus \{f\}) \cup \{e\}$. I claim T'_2 is a spanning tree. Removing f from the cycle C in $T_2 \cup \{e\}$ keeps the graph connected: any path that used f can go around the rest of C instead. So T'_2 is connected, spans all n vertices, and has $(n - 1) - 1 + 1 = n - 1$ edges. A connected graph on n vertices with $n - 1$ edges is a tree, so T'_2 is a spanning tree.

Its weight is

$$w(T'_2) = w(T_2) - w(f) + w(e) < w(T_2),$$

which contradicts the assumption that T_2 is a minimum spanning tree.

So two distinct minimum spanning trees cannot exist, i.e., the MST is unique when the edge weights are distinct.

Q6.1 Greedy knapsack filling (a)

Every feasible solution of the 0-1 problem is also a feasible solution of the fractional problem with the same value. Given a subset S with $\sum_{i \in S} w_i \leq C$, set $p_i = 1$ for $i \in S$ and $p_i = 0$ otherwise. Then every $p_i \in [0, 1]$, the total weight $\sum_i p_i w_i = \sum_{i \in S} w_i \leq C$ is feasible, and the total value $\sum_i p_i v_i = \sum_{i \in S} v_i$ is the same. In particular, an optimal 0-1 solution gives a feasible fractional solution of the same value. The fractional optimum is the maximum over a set of solutions that contains this one, so it is at least the 0-1 optimum.

Q6.2 Greedy knapsack filling (b)

No, it is not correct. Here is an input with capacity $C = 3$ and three items:

item	1	2	3
w_i	1	2	3
v_i	1	4	5

(There are no ties in value or weight, so both algorithms are fully determined.)

ValueGreedy. The largest value is $v_3 = 5$. Item 3 has weight $3 \leq C$, so the whole item fits and the capacity drops to 0. The algorithm stops with total value 5.

WeightGreedy. The smallest weight is $w_1 = 1$; take all of item 1 (value 1, capacity left 2). Next smallest is $w_2 = 2$; take all of item 2 (value 4, capacity left 0). The algorithm stops with total value $1 + 4 = 5$.

So the best of the two outputs has value 5.

A better solution. Take all of item 2 and one third of item 3: $p = (0, 1, \frac{1}{3})$. The weight is $2 + \frac{1}{3} \cdot 3 = 3 \leq C$, so it is feasible, and the value is $4 + \frac{1}{3} \cdot 5 = \frac{17}{3} > 5$. Hence the combined algorithm's output is not optimal.

(Intuitively, item 2 has the best value per unit weight, $v_2/w_2 = 2$, but it is neither the most valuable nor the lightest item, so neither greedy rule looks at it first.)

Q6.3 Greedy knapsack filling (c)

Write $\text{OPT} = q_1, \dots, q_n$ for the fractions in OPT. By the choice of j , $q_i = p_i$ for all $i < j$ and $q_j \neq p_j$. Let

$$W = \sum_{i < j} p_i w_i = \sum_{i < j} q_i w_i$$

be the weight both solutions put on the first $j - 1$ items.

Step 1: $q_j < p_j$. When the algorithm reached item j , its remaining capacity was $C - W$ (if this was 0 the algorithm had already stopped, and then $p_j = 0$). It took the largest fraction that fits, so

$$p_j = \min \left\{ 1, \frac{C-W}{w_j} \right\}.$$

OPT is feasible, so $W + q_j w_j \leq C$, i.e. $q_j \leq \frac{C-W}{w_j}$; also $q_j \leq 1$. So $q_j \leq \min \left\{ 1, \frac{C-W}{w_j} \right\} = p_j$. Since $q_j \neq p_j$, we get $q_j < p_j$.

Step 2: which fractions to exchange. Let $\delta = p_j - q_j > 0$, and let $\Delta = \delta w_j$ be the extra weight of item j we want to add. Let $L = \sum_{i > j} q_i w_i$ be the weight OPT puts on items after j , and let $F = C - W - q_j w_j - L \geq 0$ be the unused capacity of OPT. Since $p_j w_j \leq C - W$ (from Step 1),

$$\Delta = p_j w_j - q_j w_j \leq C - W - q_j w_j = L + F.$$

So there is enough room to add Δ if we use the free capacity first and then remove weight from later items. Concretely, let $\Delta' = \max \{0, \Delta - F\}$ be the weight that must be removed from items $j + 1, \dots, n$; note $0 \leq \Delta' \leq L$. Go through $i = j + 1, j + 2, \dots, n$ and remove from item i a weight x_i with $0 \leq x_i \leq q_i w_i$, taking as much as possible from each item until the removed weights add up to exactly Δ' . This is possible because $\Delta' \leq L = \sum_{i > j} q_i w_i$. Define $\text{OPT}' = q'_1, \dots, q'_n$ by

$$q'_i = q_i \ (i < j), \quad q'_j = p_j, \quad q'_i = q_i - \frac{x_i}{w_i} \ (i > j).$$

In words: add a δ -fraction of item j , and pay for it with unused capacity and with fractions of lower-density items.

Step 3: OPT' is valid. Every fraction is in $[0, 1]$: for $i < j$ nothing changed; $q'_j = p_j \in [0, 1]$; and for $i > j$, $0 \leq x_i \leq q_i w_i$ gives $0 \leq q'_i \leq q_i \leq 1$. The total weight is

$$\sum_i q'_i w_i = W + (q_j + \delta) w_j + (L - \Delta') = \underbrace{(W + q_j w_j + L)}_{= C - F} + \Delta - \Delta' = C - F + \min \{ \Delta, F \} \leq C.$$

So OPT' is a valid solution.

Step 4: OPT' is optimal. The change in value is

$$\text{val}(\text{OPT}') - \text{val}(\text{OPT}) = \delta v_j - \sum_{i > j} \frac{x_i}{w_i} v_i = \Delta r_j - \sum_{i > j} x_i r_i.$$

The items are sorted by relative value, so $r_i \leq r_j$ for every $i > j$, and every $x_i \geq 0$. Therefore

$$\sum_{i > j} x_i r_i \leq r_j \sum_{i > j} x_i = r_j \Delta' \leq r_j \Delta,$$

so $\text{val}(\text{OPT}') - \text{val}(\text{OPT}) \geq 0$. Thus OPT' is a valid solution with value at least that of OPT. Since OPT is optimal, OPT' cannot have larger value, so $\text{val}(\text{OPT}') = \text{val}(\text{OPT})$ and OPT' is optimal.


Finally, OPT' has fractions $p_1, \dots, p_{j-1}$ on the first $j - 1$ items (unchanged from OPT) and p_j on item j , as required. Repeating this exchange at each disagreement (as in Question 3(c)) gives an optimal solution equal to ALG, so RELATIVELYGREEDY is correct.

Q7.1 Greedily independent (a, optional)

Take the graph with vertices $1, \dots, 6$ and edges

$$(1, 2), (1, 3), (2, 4), (3, 4), (4, 5), (4, 6), (5, 6).$$

It is a 4-cycle $1-2-4-3-1$ glued to a triangle $4-5-6$ at vertex 4.



Run of GreedyIS. The degrees are $\deg(1) = \deg(2) = \deg(3) = \deg(5) = \deg(6) = 2$ and $\deg(4) = 4$. The smallest degree is 2, and the smallest-numbered vertex with that degree is 1. So the algorithm adds 1 to S and removes 1 and its neighbors 2, 3. The remaining graph is the triangle on 4, 5, 6, where every vertex has degree 2, so the algorithm picks 4 and removes 4, 5, 6. Nothing is left, and it returns $S = \{1, 4\}$, of size 2.

A larger independent set. $\{2, 3, 5\}$ is independent: none of $(2, 3)$, $(2, 5)$, $(3, 5)$ is an edge. It has size $3 > 2$, so the output of GREEDYIS is not a maximum independent set.

Q7.2 Greedily independent (b, optional)

(Why v_{diff} exists: ALG is a *maximal* independent set, since every vertex not in ALG was removed as a neighbor of some chosen vertex. So no independent set strictly contains ALG, and in particular ALG is not a proper subset of OPT.)

Let $d = \text{diff}$, and let G_d be the graph that remains right before the algorithm picks v_d . Two observations about G_d :

- **(O1) G_d is a forest.** It is obtained from the tree G by deleting vertices, and any subgraph of a tree has no cycles.
- **(O2) The vertices deleted before step d are exactly $v_1, \dots, v_{d-1}$ and their neighbors in G , and none of these neighbors is in OPT.** The first claim is just what the algorithm does. For the second: OPT contains $v_1, \dots, v_{d-1}$ and is independent, so it contains no neighbor of any v_i with $i < d$.

A consequence of (O2): v_d is in G_d , so it is not adjacent (in G) to any of $v_1, \dots, v_{d-1}$; otherwise it would have been deleted earlier. So every neighbor of v_d in G that is *not* in G_d is a neighbor of some v_i ($i < d$), and hence is not in OPT.

Which vertex u to exchange. By (O1), G_d is a forest, so it has a vertex of degree 0 or 1 (every tree with at least two vertices has a leaf). The algorithm picks a vertex of smallest degree, so v_d has degree 0 or 1 in G_d .

Degree 0 cannot happen. If v_d had no neighbors in G_d , then by the consequence of (O2) all its neighbors in G lie outside OPT. Then $\text{OPT} \cup \{v_d\}$ would be an independent set larger than OPT (recall $v_d \notin \text{OPT}$), contradicting that OPT is maximum.

So v_d has exactly one neighbor in G_d ; call it u . Define

$$\text{OPT}' = (\text{OPT} \setminus \{u\}) \cup \{v_d\}.$$

u is in OPT. Suppose not. The neighbors of v_d in G are u and vertices outside G_d ; the latter are not in OPT by the consequence of (O2), and $u \notin \text{OPT}$ by assumption. So v_d has no neighbor in OPT, and $\text{OPT} \cup \{v_d\}$ is an independent set of size $|\text{OPT}| + 1$, a contradiction. Hence $u \in \text{OPT}$.

OPT' is valid (independent). Take two distinct vertices $x, y \in \text{OPT}'$. If neither is v_d , then both are in OPT, so they are not adjacent. If $x = v_d$, then $y \in \text{OPT} \setminus \{u\}$. The only neighbors of v_d are u (removed) and vertices outside OPT, so y is not adjacent to v_d . So OPT' is an independent set.

OPT' is optimal. We removed $u \in \text{OPT}$ and added $v_d \notin \text{OPT}$, so $|\text{OPT}'| = |\text{OPT}|$. An independent set of the same size as a maximum independent set is itself maximum.

OPT' contains $v_1, \dots, v_d$. The vertex u lies in G_d , while $v_1, \dots, v_{d-1}$ were deleted before step d , so $u \neq v_i$ for all $i < d$. Hence $v_1, \dots, v_{d-1}$ stay in OPT', and v_d was added.

Repeating this exchange at each disagreement (exactly as in Question 3(c)) yields a maximum independent set that contains all of $v_1, \dots, v_k$. Since ALG is a maximal independent set, that set must equal ALG, so GREEDYIS returns a maximum independent set on every tree.